LEARNING BY STATE RECURRENCE DETECTION

Authors:

Bruce Rosen, James Goodwin, Jacques Vidal

Abstract:

This research investigates a new technique for unsupervised learning of nonlinear control problems. The approach is applied both to Michie

and Chambers BOXES algorithm and to Barto, Sutton and Anderson's extension, the ASE/ACE system, and has significantly improved the convergence

rate of stochastically based learning automata. Recurrence learning is a new nonlinear reward-penalty algorithm. It exploits information found during learning

trials to reinforce decisions resulting in the recurrence of nonfailing states. Recurrence learning applies positive reinforcement during the exploration of

the search space, whereas in the BOXES or ASE algorithms, only negative weight reinforcement is applied, and then only on

failure. Simulation results show that the added information from recurrence learning increases the learning rate. Our empirical results show that

recurrence learning is faster than both basic failure driven learning and failure prediction methods. Although recurrence learning has only been

tested in failure driven experiments, there are goal directed learning applications where detection of recurring oscillations may provide useful information

that reduces the learning time by applying negative, instead of positive reinforcement. Detection of cycles provides a heuristic to improve

the balance between evidence gathering and goal directed search.